

AUTOTUNING UNDER TIGHT BUDGET CONSTRAINTS: A DESIGN OF EXPERIMENTS APPROACH

Pedro Bruel, Steven Quinito Masnada, Brice Videau, Arnaud Legrand, Jean-Marc Vincent, Alfredo Goldman

April 28, 2019


Autotuning

Applying Design of Experiments to Autotuning

Results

AUTOTUNING: OPTIMIZING PROGRAM CONFIGURATIONS

Architectures for High Performance Computing



How to write **efficient code** for each of these?

Autotuning

The process of **automatically finding a configuration** of a program that optimizes an **objective**

Configurations

- Program Configuration
 - Algorithm, block size, . . .
- Source code transformation
 - Loop unrolling, tiling, rotation . . .
- Compiler configuration
 - -O2, vectorization, . . .
- . . .

Objectives


- Execution time
- Memory & power consumption
- . . .

AUTOTUNING: SEARCH SPACES

Search Spaces

Represent the **effect** of all possible **configurations** on the **objectives**

Can be difficult to explore, with multiple **local optima** and **undefined regions**



Unrolling, blocking and Mflops/s for matrix multiplication

Seymour K, You H, Dongarra J. A comparison of search heuristics for empirical code optimization. InCLUSTER 2008 Oct 1 (pp. 421-429)

Issue 1: Exponential Growth

Simple factors can generate large spaces:

- 30 boolean factors $\rightarrow 2^{30}$ combinations

Issue 2: Geometry

- Discrete or continuous factors
- “Smoothness”
- Interactions between factors

Issue 3: Measurement Time

Time to compile:

- Benchmark GPU applications: 1~10s
- Benchmark FPGA applications: 1~10min
- Industrial FPGA applications: 1~10h

AUTOTUNING: MULTIPLE APPROACHES

Popular Approaches

- Exhaustive
- Meta-Heuristics
- Machine Learning

System	Domain	Approach
ATLAS	Dense Linear Algebra	Exhaustive
INSIEME	Compiler	Genetic Algorithm
Active Harmony	Runtime	Nelder-Mead
ParamILS	Domain-Agnostic	Stochastic Local Search
OPAL	Domain-Agnostic	Direct Search
OpenTuner	Domain-Agnostic	Ensemble
MILEPOST GCC	Compiler	Machine Learning
Apollo	GPU kernels	Decision Trees

Main Issues

- These approaches **assume**:
 - A **large number of function evaluations**
 - Search space “**smoothness**”
 - Good solutions are **reachable**
- After optimizing:
 - **Learn “nothing”** about the search space
 - **Can’t explain** why optimizations work

DESIGN OF EXPERIMENTS

Factors, Levels, Experiments & Designs

- **Factors:** program parameters
- **Levels:** possible factor values
- **Experiment:** setting each factor to a level
- **Design:** a selection of experiments to run

Analysis

Experiment results can be used to:

- Identify relevant parameters
- Fit a regression model

A small design for 7 2-level factors:

Run	A	B	C	D	E	F	G
1	1	-1	1	-1	-1	1	1
2	1	1	1	-1	1	-1	-1
3	-1	1	-1	-1	1	1	1
4	-1	1	1	1	-1	1	-1
5	1	-1	-1	1	1	1	-1
6	1	1	-1	1	-1	-1	1
7	-1	-1	1	1	1	-1	1
8	-1	-1	-1	-1	-1	-1	-1

APPLYING DESIGN OF EXPERIMENTS TO AUTOTUNING

Our Approach

We are using:

- Efficient experimental designs to overcome issues related to exponential growth, geometry, and measurement time
- Analysis of variance to find relevant parameters
- User input to guide optimization

Design Requirements

- Support a large number of factors (Exponential Growth)
- Support numerical and categorical factors (Geometry)
- Minimize function evaluations (Measurement Time)

D-Optimal Designs

- Minimize variance of regression coefficient estimators
- Supports different factor types and numbers

D-OPTIMAL DESIGNS: EXAMPLE

Example

- Factors & Levels:

$$\mathbf{X} = (x_1 = (1, \dots, 5), \\ x_2 = ("A", "B", "C"))$$

- Model: $\mathbf{Y} = \mathbf{X}\beta + \varepsilon$

Source code

```
library(AlgDesign)

full_factorial <- gen.factorial(c(5, 3),
                               factors = c(2))

output <- optFederov(~ .,
                    full_factorial,
                    nTrials = 5)
```

Output

```
$D
[1] 0.5656854

$A
[1] 3.828125

$Ge
[1] 0.64

$Dea
[1] 0.57

$design
      1      5      6      10     13
X1 "-2" " 2" "-2" " 2" " 0"
X2 "1"  "1"  "2"  "2"  "3"

$rows
[1] 1  5  6 10 13
```

Linear Regression Model

The **linear regression model**:

$$y = \beta_0 + \beta_1 x_1 + \cdots + \beta_k x_k + \epsilon$$

We want to **estimate** $\beta_{0,\dots,k}$:

- Using $n > k$ **observations** $y_{1,\dots,n}$
- **Distinct** x_{i1}, \dots, x_{ik} , $i = 1, \dots, n$

Experiments represented by:

$$y_i = \beta_0 + \beta_1 x_{i1} + \cdots + \beta_k x_{ik} + \epsilon_i$$

Ordinary Least Squares Estimator $\hat{\beta}$

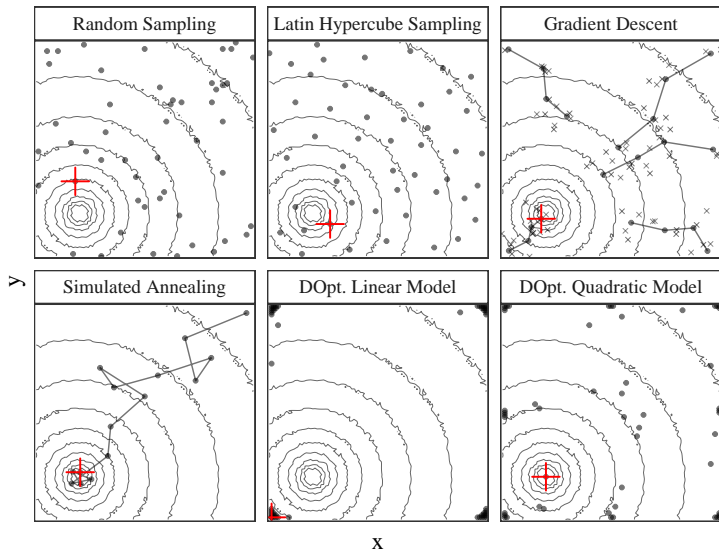
$$\hat{\beta} = (X^T X)^{-1} X^T Y$$

The **variance** of $\hat{\beta}$ is proportional to the **covariance matrix** $(X^T X)^{-1}$

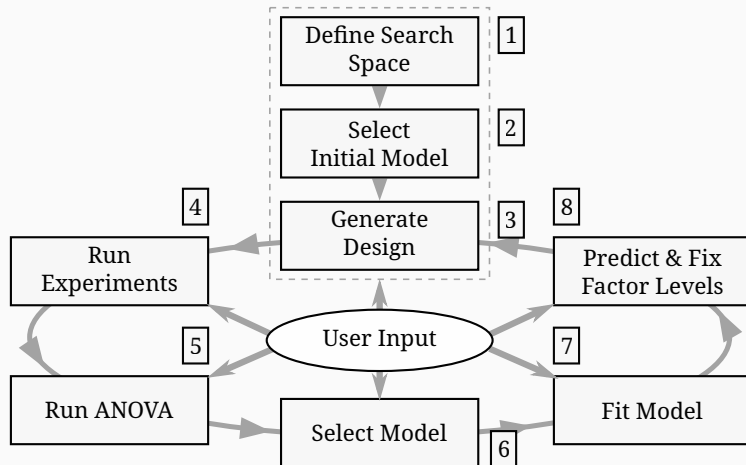
Design Criteria using $(X^T X)^{-1}$

- **D: determinant**, minimizes generalized variance of $\hat{\beta}$
- **A: trace**, average variance of $\hat{\beta}$

SAMPLING STRATEGIES



A DESIGN OF EXPERIMENTS APPROACH TO AUTOTUNING



GPU LAPLACIAN KERNEL: A MOTIVATING EXAMPLE


The Search Problem

- Relatively **small valid search space**
- **Completely evaluated**
- Known **global optimum**
- Known **model approximation**
- **Budget of 125 points**

Initial Model

$$\begin{aligned} \text{cost} = & y_component_number + 1/y_component_number + \\ & vector_length + lws_y + 1/lws_y + \\ & load_overlap + temporary_size + \\ & elements_number + 1/elements_number + \\ & threads_number + 1/threads_number \end{aligned}$$

Results



GPU LAPLACIAN KERNEL: A MOTIVATING EXAMPLE

	Mean	Max
RS	120.00	125.00
LHS	98.92	125.00
GS	22.17	106.00
GSR	120.00	120.00
GA	120.00	120.00
LM	119.00	119.00
DLMT	54.84	56.00

Table 1: Points used by applications

Summary

Our approach:

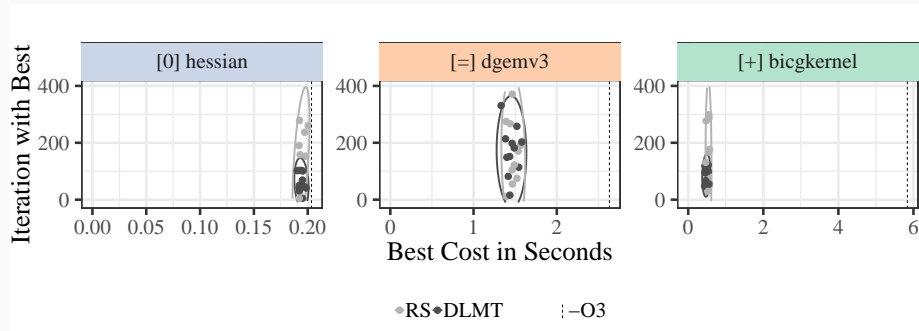
- Was **always close to the optimum**
- Used **half of the budget**

SPAPT: SEARCH PROBLEMS IN AUTOMATIC PERFORMANCE TUNING


Kernel	Operation	Factors	Size
atax	Matrix transp. & vector mult.	18	2.6×10^{16}
dgemv3	Scalar, vector & matrix mult.	49	3.8×10^{36}
gemver	Vector mult. & matrix add.	24	2.6×10^{22}
gesummv	Scalar, vector, & matrix mult.	11	5.3×10^9
hessian	Hessian computation	9	3.7×10^7
mm	Matrix multiplication	13	1.2×10^{12}
mvt	Matrix vector product & transp.	12	1.1×10^9
tensor	Tensor matrix mult.	20	1.2×10^{19}
trmm	Triangular matrix operations	25	3.7×10^{23}
bicg	Subkernel of BiCGStab	13	3.2×10^{11}
lu	LU decomposition	14	9.6×10^{12}
adi	Matrix sub., mult., & div.	20	6.0×10^{15}
jacobi	1-D Jacobi computation	11	5.3×10^9
seidel	Matrix factorization	15	1.3×10^{14}
stencil3d	3-D stencil computation	29	9.7×10^{27}
correlation	Correlation computation	21	4.5×10^{17}

Balaprakash P, Wild SM, Norris B. SPAPT: Search problems in automatic performance tuning. Procedia Comp. Sci. 2012 Jan 1;9:1959-68.

SPAPT: PRELIMINARY RESULTS



SPAPT: PRELIMINARY RESULTS



Experimental Settings

- Using the same model for all applications
- Fixed number of iterations
- Automated approach

Summary

- Performance similar to random sampling
- Using less points

CONCLUSION

Summary

Our approach uses:

- Efficient experimental designs to overcome issues related to exponential growth, geometry, and measurement time
- Analysis of variance to find relevant parameters
- User input to guide optimization

Perspectives


- Explore tailored models for each application
- Leverage user input and analysis
- Use our approach to autotune industrial-level FPGA applications

AUTOTUNING UNDER TIGHT BUDGET CONSTRAINTS: A DESIGN OF EXPERIMENTS APPROACH

Pedro Bruel, Steven Quinito Masnada, Brice Videau, Arnaud Legrand, Jean-Marc Vincent, Alfredo Goldman

April 28, 2019

SPAPT: PRELIMINARY RESULTS



DESIGN EFFICIENCY: INTRODUCTION

Linear Regression Model

A simple **regression model**:

$$y = \beta_0 + \beta_1 x_1 + \cdots + \beta_k x_k + \epsilon$$

We want to **estimate** $\beta_{0,\dots,k}$:

- Using $n > k$ **observations** $y_{1,\dots,n}$
- **Distinct** x_{i1}, \dots, x_{ik} , $i = 1, \dots, n$

We will use n **experiments** such as:

$$y_i = \beta_0 + \beta_1 x_{i1} + \cdots + \beta_k x_{ik} + \epsilon_i$$

Least Squares Method

Writing in **matrix form** we get:

$$Y = X\beta + \epsilon$$

The **least squares method** aims to minimize:

$$\begin{aligned} L &= \sum_{i=1}^n \epsilon_i^2 = \epsilon^T \epsilon = (Y - X\beta)^T (Y - X\beta) = \\ &= Y^T Y - \beta^T X^T Y - Y^T X \beta + \beta^T X^T X \beta = \\ &= Y^T Y - 2\beta^T X^T Y + \beta^T X^T X \beta \end{aligned}$$

DESIGN EFFICIENCY: ESTIMATING MODEL COEFFICIENTS

Minimizing Least Squares

The **least squares method** aims to minimize:

$$L = Y^T Y - 2\beta^T X^T Y + \beta^T X^T X \beta$$

Derivative with respect to β , **evaluated** at $\hat{\beta}$:

$$\left. \frac{\partial L}{\partial \beta} \right|_{\hat{\beta}} = -2X^T Y + 2X^T X \hat{\beta} = 0$$

Where $\hat{\beta}$ is an **estimator** of β

Computing $\hat{\beta}$

The previous equation simplifies to:

$$\hat{\beta} = (X^T X)^{-1} X^T Y$$

The estimator $\hat{\beta}$ is proportional to $(X^T X)^{-1}$

Dispersion or Covariance Matrix

- **Information matrix:** $X^T X$
- **Dispersion or Covariance matrix:** $(X^T X)^{-1}$

DESIGN EFFICIENCY: THE DISPERSION MATRIX

Computing $(X^T X)^{-1}$

A design $D_{n,2}$, with **2-level factors**, will have a 3×3 **dispersion matrix**, if we assume **linear relationships** and no **factor interactions**:


```
factorial <- gen.factorial(c(2, 2))  
model <- model.matrix(~., factorial)  
dispersion <- t(model) %*% model  
eigen(dispersion)$values
```

	(Intercept)	X1	X2
(Intercept)	4	0	0
X1	0	4	0
X2	0	0	4

```
[1] 4 4 4
```

Interpreting Eigenvalues of $(X^T X)^{-1}$

The **eigenvalues** $\lambda_{0,1,2}$ of the **dispersion matrix** can represent its “size”:



We can **minimize the coefficient estimator** $\hat{\beta}$ by **minimizing the eigenvalues** of $(X^T X)^{-1}$

DESIGN EFFICIENCY: METRICS

Defining a Design

Consider a design $D_{n,k-1}$:

- $x_{1,\dots,k-1}$ 2-level factors
- n experiments

Its $k \times k$ dispersion matrix $(X^T X)^{-1}$:

- Constructed using the linear model:
 - $Y = \beta X + \epsilon$
- With eigenvalues λ_0, \dots, m

We can define efficiency metrics for β based on the eigenvalues of the dispersion matrix

Some Efficiency Metrics based on $(X^T X)^{-1}$

A-Efficiency

$$A_{eff} = \left(n \times \text{tr} \left((X^T X)^{-1} \right) / k \right)^{-1}, A_{eff} \in [0, 1]$$


“Arithmetic mean” of eigenvalues of $(X^T X)^{-1}$

D-Efficiency

$$D_{eff} = \left(n \times \left| (X^T X)^{-1} \right|^{1/k} \right)^{-1}, D_{eff} \in [0, 1]$$

“Geometric mean” of eigenvalues of $(X^T X)^{-1}$

AUTOTUNING: SEARCH SPACES



Compiler impact on performance

Seymour K, You H, Dongarra J. A comparison of search heuristics for empirical code optimization. InCLUSTER 2008 Oct 1 (pp. 421-429)

Our Approach

Using **efficient experimental design** to overcome issues related to **exponential growth**, **geometry**, and **measurement time**

Design Requirements

- Support a large number of factors
(**Exponential Growth**)
- Support numerical and categorical factors
(**Geometry**)
- Minimize function evaluations
(**Measurement Time**)

Main Design Candidates

Screening designs:

- Estimate **main effects**
- Aim to **minimize runs**
- Assume **interactions are negligible**

Mixed-Level designs:

- Factors have **different numbers of levels**
- Many **optimality criteria**

SCREENING DESIGNS

A Plackett-Burman screening design for 7
2-level factors:

Run	A	B	C	D	E	F	G
1	1	-1	1	-1	-1	1	1
2	1	1	1	-1	1	-1	-1
3	-1	1	-1	-1	1	1	1
4	-1	1	1	1	-1	1	-1
5	1	-1	-1	1	1	1	-1
6	1	1	-1	1	-1	-1	1
7	-1	-1	1	1	1	-1	1
8	-1	-1	-1	-1	-1	-1	-1

Screening Designs

Plackett-Burman designs for 2-level factors:

- Orthogonal arrays of strength 2
- Estimate the main effects of n factors with $n + 1$ runs


Construction:

- For $n + 1$ multiple of 4
- Identical to a fractional design if $n + 1$ is a power of two

LOOKING AT DATA: CUDA COMPILER FLAGS

CUDA Compiler Flags

- Rodinia benchmark
- 15 factors, few with multiple levels
- 10^6 combinations
- 1~10s to measure
- Screening experiment:
 - 15 “2-level” factors
 - 4 “dummy” factors



MIXED-LEVEL DESIGNS

A multi-level design for 1 2-level factor and 3 3-level factors:

Run	A	B	C	D
1	1	1	1	3
2	1	1	2	1
3	1	1	3	2
4	1	2	1	2
5	1	2	2	3
6	1	2	3	1
7	1	3	1	1
8	1	3	2	2
9	1	3	3	3
10	2	1	1	1
11	2	1	2	2
12	2	1	3	3
13	2	2	1	3
14	2	2	2	1
15	2	2	3	2
16	2	3	1	2
17	2	3	2	3
18	2	3	3	1

Mixed-Level Designs

Strategy 1: Contractive Replacement

- Find specific sets of k -level columns of a design, contract the set into a new factor of with more levels
- Maintain orthogonality of the design

Strategy 2: Direct Construction

Directly generate small mixed-level designs by solving Mixed Integer Programming problems

Strategy 3: D-Optimal Designs

LOOKING AT DATA: FPGA COMPILER PARAMETERS

FPGA Compiler Parameters

- CHStone benchmark
- 141 factors, most with multiple levels
- 10^{128} combinations
- 1~10min to measure
- Multiple objectives
- Search with meta-heuristics:
 - Unstructured data difficults analysis
 - We are working on obtaining more data

